

Communication-Efficient Approaches to Federated Deep Neural Networks

Bachelor's Thesis

By: Adrian Edward Thomas Henkel

adrian.henkel@campus.lmu.de


Supervisor: Reza Nasirigerdeh

reza.nasirigerdeh@tum.de

Advisors: Prof. Dr. Jan Baumbach, and Dr. Josch Pauling

jan.baumbach@uni-hamburg.de

josch.pauling@wzw.tum.de



Agenda

- Introduction
 - Deep neural networks
 - Federated learning
- Communication-efficient approaches
- Simulation framework
- Datasets and models
- Results
- Summary and outlook


A typical neural network

Architecture:

- One Input layer
- One or multiple hidden layers
- One output layer

Training process:


- Initiate the model with random weights
- Iteratively update the network to minimize a loss function
- In each iteration:
 - Select a subset of the training data (batch)
 - Find updated weights that optimize the loss function of the batch (backpropagation)



Federated learning

- **Multiple clients** train a **global model** under the coordination of a **central server**

1. Each **client** trains the model on its local data
2. Each **client** shares the updated model with the server
3. The **server** computes the global model by taking the weighted average over the updated local models from the clients
4. The **server** sends the global model back to the clients
5. Repeat step 1-4 until the global model converges



- Privacy
 - Reconstruction of the private data from the model parameters is possible in FL
- Network communication
 - A huge amount of traffic might be exchanged over the network in FL
- Heterogeneous configurations
 - Clients with various computational and communication speeds
 - Non-IID (Independent and Identically Distributed) data across clients
- In this work, we **focus** on the **network communication challenge in deep neural networks**.

FL - Network bandwidth usage

K = number of clients
 G = number of model parameters
 L = size of each model parameter in bits
 N = number of iterations (communication rounds)
 π_C = network bandwidth usage

$$\rightarrow \pi_C = 2 \cdot K \cdot G \cdot L \cdot N$$

- Communication-efficient approaches
 - Gradient quantification (↓ L)
 - Gradient sparsification (↓ G)
 - More local updates (↓ N)


Example:

- $K = 50$
 - $G = 1.000.000$
 - $L = 32$
 - $N = 200$
- $\approx 80\text{GB}$**

Gradient quantification (GQ)

- Default size of a parameter is 32 bits ($L=32$)
- Reduce the size of each parameter to 16 bits before sending the model
- The gradients are re-transformed into 32-bit representation before training


$$\pi_{GQ} = 2 \cdot K \cdot G \cdot \frac{L}{2} \cdot N = \frac{\pi_C}{2}$$



Gradient sparsification (GS)

- A model might be over-parameterized
- Calculate the difference between the global model and the updated model
- Eliminate all parameters under a Percentile P and determine their positions from a binary matrix
- Send the sparse model and the binary matrix

$$\pi_{GS}(P) = \left(\frac{\sum_{i=1}^N \frac{F_i^{server}}{100}}{2N} + \frac{\sum_{j=1}^K \sum_{i=1}^N \frac{F_i^j}{100}}{2K \cdot N} + \frac{1}{L} \right) \cdot \pi_C$$



Parameter after training


Parameter over the percentile

Old parameters of the global model

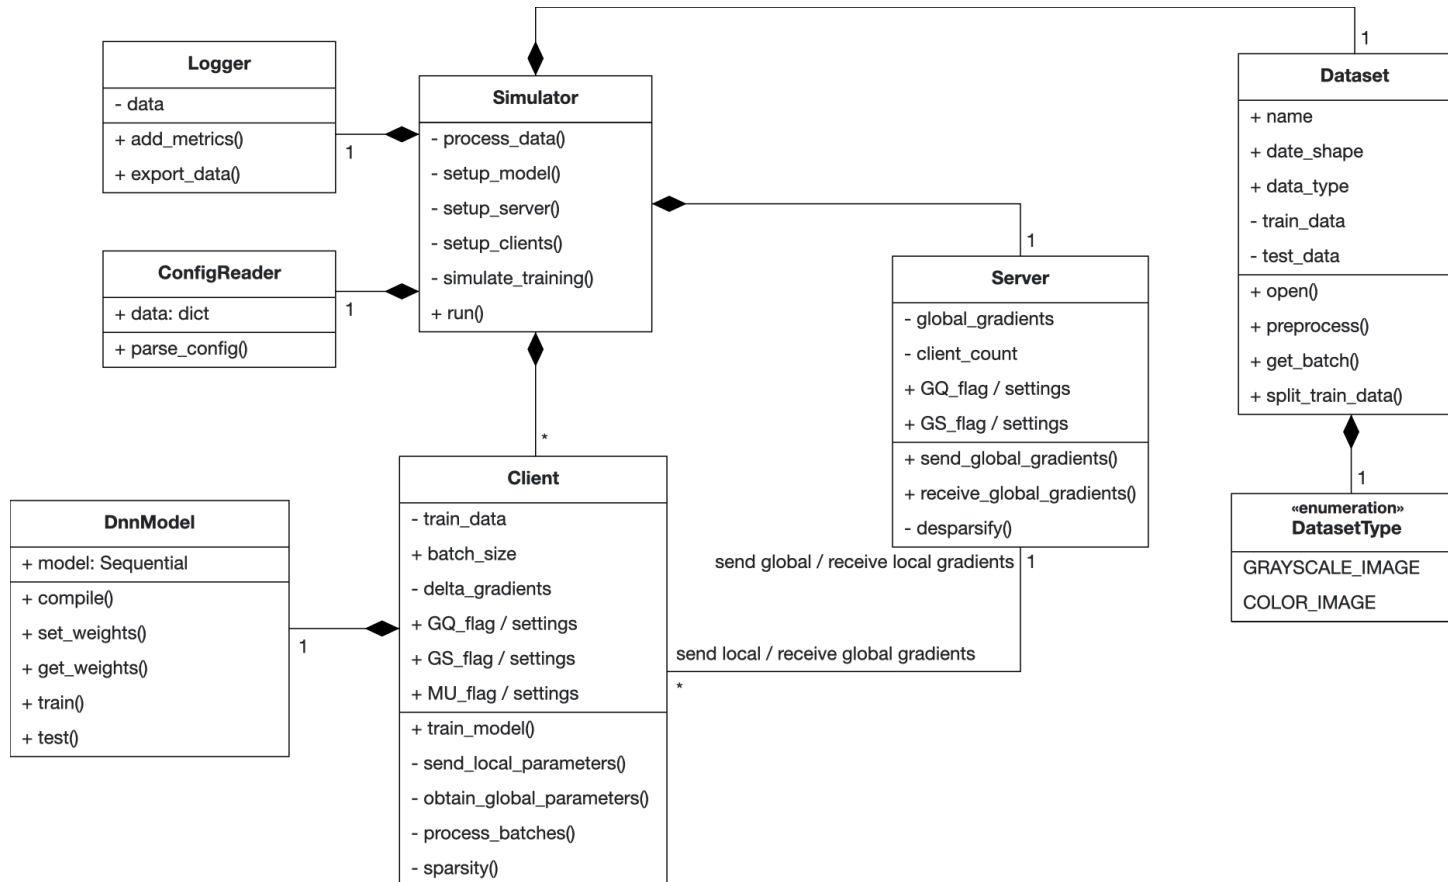
More local updates (MU)

- Typically the number of local updates (E) is one
- Increase of the local updates in one iteration
- Faster convergence of the global model
-> less iterations

$$\pi_{MU} = 2 \cdot K \cdot G \cdot L \cdot N' = \frac{N'}{N} \cdot \pi_C$$



Simulation framework




Repository: <https://gitlab.lrz.de/000000000149C8EB/com-eff>



NumPy

Datasets





Fashion-MNIST:

- Created 2017
- Zalando assortment
- 70000 samples
- 28 x 28 px. grayscale
- 10 classes

CIFAR-10

- Created 2009
- Subset of tiny image DB
- 60000 samples
- 32 x 32 px. RGB
- 10 classes



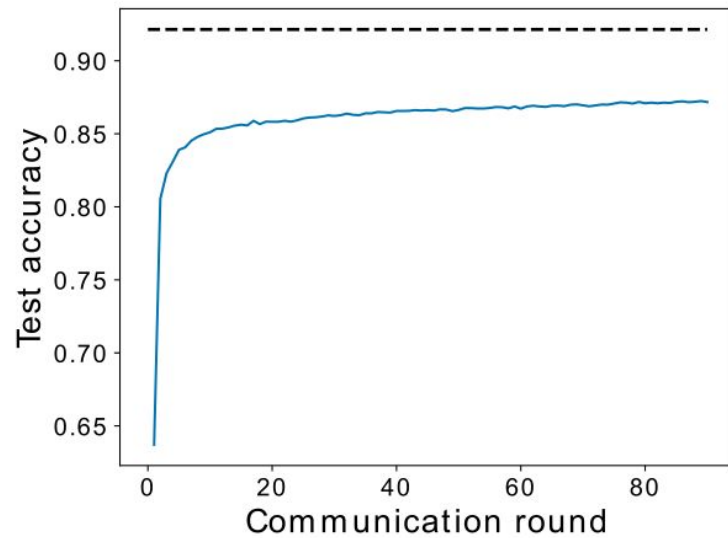
Colorectal Histology Dataset:

- Created 2009
- 4000 samples
- 150 x 150 px. RGB
- 8 classes


- **2CFNN**
 - Two convolutional layers
 - Two Max-Pooling layers
 - Dense layer with 512 neurons as penultimate layer
 - Total of 1,633,370 trainable parameters
- **3CFNN**
 - Three convolutional layers
 - Two Max-Pooling layers
 - Dense layer with 1024 neurons as penultimate layer
 - Total of 9,878,794 trainable parameters
- **VGG16**
 - Five blocks of two or three convolutional layers
 - Convolutional layers are ranging from 64 to 512 filters
 - Total of 65,087,304

Results - Accuracy


Gradient quantification



2CFNN - FMNIST




3CFNN - CIFAR-10




VGG16 - CCH

Results - Accuracy


Gradient sparsification



2CFNN - FMNIST






3CFNN - CIFAR-10



VGG16 - CCH

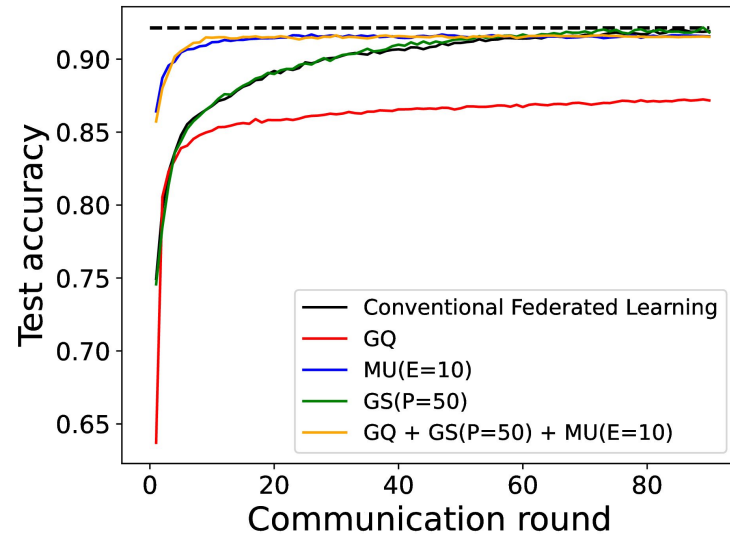
Results - Accuracy

More local updates




Results - Accuracy

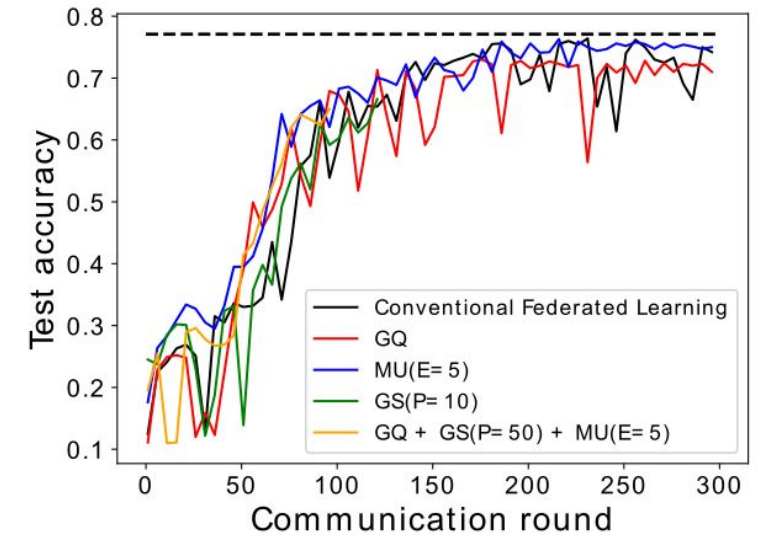
Combinations



2CFNN - FMNIST



3CFNN - CIFAR-10



VGG16 - CCH

Results - Bandwidth usage

| Target accuracy \ Approach | 0.84 | 0.86 | 0.88 | 0.90 | 0.92 |
|----------------------------|-------------------------|----------------------------|---------------------------|----------------------------|-----------------------------|
| P=10 | 665.35 642.48 3.44 | 1064.56 1156.46 -8.63 | 1996.04 1927.43 3.44 | 3725.95 4240.35 -13.81 | 11177.85 11179.10 -0.01 |
| P=50 | 665.35 509.41 23.44 | 1064.56 815.05 23.44 | 1996.04 1528.22 23.44 | 3725.95 2852.68 23.44 | 11177.85 7437.35 33.46 |
| P=90 | 665.35 376.34 43.44 | 1064.56 677.41 36.37 | 1996.04 1053.75 47.21 | 3725.95 1655.89 55.56 | 11177.85 3989.18 64.31 |
| E=2 | 665.35 399.21 40.00 | 1064.56 532.28 50.00 | 1996.04 1064.56 46.67 | 3725.95 1996.04 46.43 | 11177.85 7318.83 34.52 |
| E=10 | 665.35 133.07 80.00 | 1064.56 133.07 87.50 | 1996.04 266.14 86.67 | 3725.95 665.35 82.14 | - |
| E=20 | 665.35 133.07 80.00 | 1064.56 133.07 87.50 | 1996.04 266.14 86.67 | 3725.95 399.21 89.29 | 11177.85 9181.80 17.86 |
| GQ | 665.35 399.21 40.00 | 1064.56 1663.37 -56.25 | - | - | - |
| GQ + GS(P=50) + MU(E=10) | 665.35 50.01 92.48 | 1064.56 100.01 90.61 | 1996.04 150.02 92.48 | 3725.95 200.02 94.63 | - |

Table 5.1: *2CFNN-FMNIST*: Bandwidth usage by the conventional federated training (in MB) | Bandwidth usage by the approach (in MB) | Bandwidth saving of the approach (in percent) to reach the target accuracy

Results - Bandwidth usage

| Target accuracy \ Approach | 0.5 | | | 0.6 | | | 0.68 | | | 0.7 | | | 0.72 | | | 0.733 | | |
|----------------------------|---------|---------|-------|----------|----------|-------|----------|----------|-------|----------|----------|-------|----------|----------|-------|----------|----------|-------|
| P=10 | 7112.73 | 6868.26 | 3.44 | 13435.16 | 12973.36 | 3.44 | 23709.11 | 20604.73 | 13.09 | 26870.32 | 24420.42 | 9.12 | 36353.96 | 30525.52 | 16.03 | 60063.07 | 54182.78 | 9.79 |
| P=50 | 7112.73 | 5445.69 | 23.44 | 13435.16 | 9076.15 | 32.44 | 23709.11 | 15126.91 | 36.20 | 26870.32 | 17547.22 | 34.70 | 36353.96 | 19967.52 | 45.07 | 60063.07 | 25413.21 | 57.69 |
| P=90 | 7112.73 | 6655.84 | 6.42 | 13435.16 | 10891.38 | 18.93 | 23709.11 | 16337.06 | 31.09 | 26870.32 | 18757.37 | 30.19 | - | - | - | - | - | - |
| E=2 | 7112.73 | 4741.82 | 33.33 | 13435.16 | 9483.64 | 29.41 | 23709.11 | 15806.07 | 33.33 | 26870.32 | 18967.28 | 29.41 | 36353.96 | 26080.02 | 28.26 | - | - | - |
| E=10 | 7112.73 | 1580.61 | 77.78 | 13435.16 | 2370.91 | 82.35 | 23709.11 | 6322.43 | 73.33 | 26870.32 | 14225.46 | 47.06 | - | - | - | - | - | - |
| E=20 | 7112.73 | 790.30 | 88.89 | 13435.16 | 1580.61 | 88.24 | 23709.11 | 11854.55 | 50.00 | - | - | - | - | - | - | - | - | - |
| GQ | 7112.73 | 3556.37 | 50.00 | 13435.16 | 6717.58 | 50.00 | 23709.11 | 11064.25 | 53.33 | 26870.32 | 13040.01 | 51.47 | 36353.96 | 16201.22 | 55.43 | 60063.07 | 19757.59 | 67.11 |
| GQ + GS(P=50) + MU(E=2) | 7112.73 | 1781.89 | 74.95 | 13435.16 | 3563.78 | 73.47 | 23709.11 | 5939.63 | 74.95 | 26870.32 | 6830.57 | 74.58 | 36353.96 | 10097.37 | 72.22 | - | - | - |

Table 5.2: *3CFNN-CIFAR-10*: Bandwidth usage by the conventional federated training (in MB) | Bandwidth usage by the approach (in MB) | Bandwidth saving of the approach (in percent) to reach the target accuracy

Results - Bandwidth usage

| Target accuracy \ Approach | 0.55 | 0.65 | 0.7 | 0.72 | 0.76 | 0.771 |
|----------------------------|-------------------------|--------------------------|-------------------------|-------------------------|-------------------------------|-------------------------|
| P=10 | 205.68 202.92 1.34 | 236.92 262.94 -10.98 | - | - | - | - |
| P=90 | 205.68 164.93 19.81 | - | - | - | - | - |
| E=2 | 205.68 145.80 29.11 | 236.92 221.30 6.59 | 343.66 275.97 19.70 | 367.09 328.04 10.64 | 466.03 377.506 36 18.95 | 643.06 387.92 39.68 |
| E=5 | 205.68 166.62 18.99 | 236.92 221.30 6.59 | 343.66 294.20 14.39 | 367.09 333.25 9.22 | 466.03 557.15 -19.55 | 643.06 577.98 10.12 |
| E=10 | 205.68 46.86 77.22 | 236.92 78.10 67.03 | 343.66 91.12 73.48 | 367.09 124.97 65.96 | - | - |
| E=20 | 205.68 41.66 79.75 | 236.92 67.70 71.43 | 343.66 119.76 65.15 | 367.09 130.17 64.54 | - | - |
| GQ | 205.68 91.12 55.70 | 236.92 100.23 57.69 | 343.66 135.38 60.61 | 367.09 184.85 49.65 | - | - |
| GQ + GS(P=50) + MU(E=5) | 205.68 76.06 63.02 | 236.92 102.53 56.72 | - | - | - | - |

Table 5.3: *VGG16-CCH*: Bandwidth usage by the conventional federated training (in GB) | Bandwidth usage by the approach (in GB) | Bandwidth saving of the approach (in percent) to reach the target accuracy

Summary

- **All three approaches** can **significantly save** the network **bandwidth** but also affect the accuracy of the models
- The **More local update** approach is the **most suitable** choice taking communication **efficiency and accuracy** into account
- The **More local update** approach is the **most contributing** approach in the **combination**
- **Combining** the approaches **provides further bandwidth savings** in comparison with each individual approach

- Gradient sparsification
 - Multi-threaded implementation (one thread per layer) to improve the sparsification speed
 - Server-side sparsification
- Communication-efficient approaches in
 - Non-IID label distributions
 - Imbalanced sample size distributions